

RiskPulse: AI/NLP Financial Risk Intelligence & Strategic Portfolio Stress Testing Platform


Interpretable NLP Risk Scoring (-1 to +1) • 8-Class Event Taxonomy • Dynamic Module B Stress Testing

CANDIDATE & INSTITUTION DETAILS	PROJECT SPECIFICATIONS & HIGHLIGHTS
• Candidate: Abhi Pandey • College: VIT Bhopal University • Degree: B.Tech Computer Science and Engineering (Artificial Intelligence and Machine Learning) • Hackathon: S&P; Global & CRISIL Campus Hackathon 2026 • Problem Statement: AI/NLP Financial Risk Intelligence & Module B Stress Testing	• Dual-Source Ingestion: Financial News Wire + Social/Market Feed • Live RSS Ingestion: Public financial feed with offline fallback • NLP Risk Engine: Deterministic scoring [-1.0 to +1.0], 8 event classes, 1-10 impact • Module B Stress Testing: Activated automatically when Impact ≥ 7 • Automated Tests: 30/30 unit tests passed (100% pass rate) • Architecture: React 18 + FastAPI + SQLite with sub-15ms latency

Problem Statement & Solution Approach

THE PROBLEM: LATENCY & BLIND SPOTS	THE SOLUTION: DETERMINISTIC RISK INTELLIGENCE
• Lagging Risk Indicators: Traditional institutional credit analysis relies heavily on quarterly SEC filings (10-K/10-Q) and periodic rating agency updates, creating massive latency. • Unstructured Data Noise: Emergent market shocks erupt first across news wires, regulatory whispers, and social channels, but are too unstructured for legacy systems. • Disconnected Valuation Shocks: Sentiment indicators rarely translate directly into balance sheet haircuts or quantified portfolio drawdowns. • Black-Box Hallucination: Opaque LLMs present compliance and governance risks under Basel III due to stochastic outputs and lack of explainability.	• Dual-Source Ingestion: Normalizes structured news wires and unstructured social feeds, with optional live RSS feeds and offline local resilience. • Interpretable NLP Risk Engine: Lexicon-calibrated sentiment [-1.0 to +1.0], 8-class event classifier, and 1-10 impact scoring formula with 100% deterministic reproducibility. • Downstream Module B Integration: Automatically triggers stress testing when Impact Score ≥ 7, applying sector-specific haircuts across equities, tech, and fixed income. • Institutional Transparency: Complete before-and-after mark-to-market valuation, dollar loss, and drawdown percentage visible on an interactive React dashboard.

End-to-End System Architecture



Technical Implementation & Tech Stack

BACKEND & REST API	INTERPRETABLE NLP RISK ENGINE
• Core Framework: FastAPI (Python 3.10+) with Uvicorn ASGI server • Storage: SQLite relational database with dynamic path resolution (/tmp for serverless Vercel) • Validation: Strict Pydantic v2 schemas for all inputs and outputs • Endpoints: GET /health, POST /analyze, GET /signals, POST /ingest, GET /portfolio, POST /stress-test	• Sentiment Analysis: Financial sentiment lexicon normalized from -1.00 to +1.00 • Event Taxonomy: 8 distinct categories: Regulatory, Geopolitical, Earnings, Operational, Cyber, Supply Chain, M&A, Macroeconomic • Impact Formulator: $\text{Impact} = \min(10, \text{round}(\text{Sentiment} *5 + \text{EventWeight}*3 + \text{Conf}*2))$ • Latency: Sub-15ms deterministic inference with zero GPU dependency
MODULE B: PORTFOLIO STRESS TESTING	FRONTEND DASHBOARD & TEST SUITE
• Baseline Portfolio: \$1,000,000 synthetic multi-asset portfolio (Equities, Tech, Fixed Income, Crypto) • Trigger Condition: Impact Score ≥ 7 triggers automated scenario analysis • Asset Haircuts: Sector shocks (e.g. Regulatory: Tech -15%, Equities -8%, Crypto -25%) • Risk Metrics: Pre-stress, Post-stress valuation, absolute loss (\$), and drawdown (%)	• Frontend: React 18 SPA with Vite, Tailwind CSS, and Lucide React icons • Resilience: Built-in benchmark fallback data guarantees zero cold-start breakage • Live RSS: Interactive header control triggers live financial feeds with graceful fallback • Verification: 30/30 unit and integration tests passing in automated CI/CD suite

Key Measured Results

30 / 30 Automated Tests Passed	< 15 ms Avg NLP Latency	-16.4 % Max Portfolio Drawdown	100 % Explainable Auditability
--	--------------------------------------	--	--

Event / Target Entity	Source	Sentiment	Event Category	Impact	Risk Level	Stress Trigger	Drawdown %
NVIDIA DOJ Subpoena	financial_news	-0.84	Regulatory	8 / 10	High	ACTIVATED	-16.4 %
Tesla Autopilot Probe	social_feed	-0.76	Regulatory	8 / 10	High	ACTIVATED	-14.2 %
CrowdStrike Cloud Outage	financial_news	-0.92	Cyber Attack	9 / 10	Severe	ACTIVATED	-15.8 %
Taiwan Strait Logistics Halt	financial_news	-0.71	Supply Chain	7 / 10	High	ACTIVATED	-11.5 %
JPMorgan Q3 Earnings Beat	financial_news	+0.81	Earnings	3 / 10	Low	Bypassed (<7)	No Stress
Apple Product Launch Buzz	social_feed	+0.68	Operational	2 / 10	Low	Bypassed (<7)	No Stress

Business Value & Financial Domain Impact

1. EARLY WARNING SURVEILLANCE	2. REGULATORY & AUDIT COMPLIANCE
Detects emergent counterparty credit deterioration, geopolitical escalations, and supply chain disruptions hours before traditional credit rating agencies publish rating watches or revisions. Synthesizes news wire velocity and market social buzz into immediate signals.	Eliminates generative LLM hallucination risk. Every sentiment score, event classification, and impact calculation is mathematically derived from transparent, auditable lexical tokens—meeting Basel III and CRISIL governance standards.
3. DIRECT VALUATION & STRESS LINKAGE	4. ZERO VENDOR LOCK-IN & CLOUD COST
Bridges the gap between qualitative financial journalism and quantitative risk management. Automates conditional execution (Impact Score ≥ 7) to estimate mark-to-market loss and portfolio value at risk across multiple asset classes.	Engineered entirely with open-source Python, SQLite, and React. Operates fully air-gapped without expensive recurring terminal subscriptions (Bloomberg/Refinitiv), enabling immediate deployment on any serverless or local institutional infrastructure.

Limitations & Strategic Roadmap

CURRENT SYSTEM LIMITATIONS	STRATEGIC ROADMAP & ENHANCEMENTS
• Lexicon Boundary: While deterministic and explainable, rule-based dictionary matching lacks deep syntactic parsing for subtle sarcasm, double negations, and complex multi-clause clauses. • Synthetic Benchmark Portfolio: Stress tests currently evaluate a normalized \$1,000,000 multi-asset synthetic benchmark rather than live custodian accounts. • Asset Class Haircut Granularity: Shocks are applied at sector and asset-class levels rather than via continuous multi-factor regression models. • Public RSS Volatility: Free public financial RSS feeds are subject to third-party network throttling; isolated with fallback for resilience.	• Hybrid FinBERT Embeddings: Integrate quantized domain-specific transformers (FinBERT / Llama-3 8B) with deterministic guardrails to capture semantic nuance while preserving explainability. • Live Custodian API Integration: Connect FIX / REST broker protocols (Interactive Brokers, Alpaca, Bloomberg EMSX) for real-time dynamic portfolio rebalancing. • Historical Crisis Backtesting: Validate shock formulas against empirical historical crises (2008 Lehman collapse, 2020 COVID crash, 2023 SVB banking shock). • Monte Carlo & Expected Shortfall: Expand Module B with parametric 99% VaR and Expected Shortfall under adverse macroprudential regulatory mandates.